

LUIS ZERMENO

AI Application Engineer

Berlin, Germany · lgzg90@hotmail.com · +49 1512 9712847 ·
lugz.dev · github.com/louiszcode · [LinkedIn](#)

PROFILE

AI Application Engineer building and shipping LLM-powered products end to end: RAG pipelines, multi-agent systems, async APIs, and production deployments. Three live products with paying users, including a multi-LLM consensus engine for financial analysis and a production SaaS for YouTube transcription and translation. Four years in B2B SaaS talent acquisition across DACH/EMEA before the pivot, which means I already know how software teams ship and how to talk to the humans around them.

SELECTED PROJECTS

AgoraFinancials — Multi-LLM financial analysis platform · agorafinancials.com

- Built the Pythagoras Method: 9 LLMs across 5 providers independently rate 8 financial metrics, with structured multi-round debate loops that resolve disagreements until consensus. Functions as a homegrown eval-and-consensus harness to reduce single-model bias.
- Two-tier model architecture (fast models for numerical extraction, deep models with RAG tools for qualitative reasoning over SEC filings) behind a unified OpenRouter gateway.
- RAG pipeline on SEC filings with pgvector (chosen over Pinecone to keep infra single-database); LangChain create_agent with retrieval tools; sliding-window debate compression to keep token cost flat across rounds.
- Production: FastAPI, PostgreSQL + pgvector, Stripe token economy, PDF report generation, Dockerized on Railway.

TubeText — YouTube transcription + translation SaaS · tubetext.app

- Async FastAPI backend orchestrating Deepgram Nova-3 (audio→text), OpenAI GPT-4 mini via LangChain (structured summaries), and Cerebras Llama-3.3-70b (real-time translation streamed over SSE) in four languages.
- YAML-driven prompt management for maintainability; async SQLAlchemy with Alembic migrations; Google OAuth and Stripe billing; PDF export.
- Chose Deepgram over Whisper for production reliability; FastAPI for async-first performance; Dockerized and deployed on Railway.

Spralingua — Voice-first AI language learning platform · spralingua.com

- Multi-agent architecture with specialized agents for vocabulary, grammar, and conversation, built on Claude's language understanding. Sub-2-second conversation loop end to end.
- Full-stack Flask app with real-time speech-to-text (Web Speech API), Minimax TTS for natural multilingual speech, session-based learning progression, and PostgreSQL persistence.
- Serving users in 4 languages (Spanish, German, Portuguese, English). Deployed on Railway with Unicorn.

TECHNICAL EXPERIENCE

AI / Automation Engineering — in-role at Lucanet AG

Berlin (Remote) · October 2023 – February 2026

- Built an internal AI recruitment assistant in Python with LangChain and the Claude API: automated candidate screening, structured interview-feedback extraction from transcriptions, and communications across high-volume GTM pipelines.
- Led ATS migration from Personio to Greenhouse — scoped data model, managed data transfer, drove team adoption; standardized hiring process for a DACH/EMEA-wide team.

Internal Tooling — in-role at RapidAPI

Berlin · June 2022 – August 2023

- Designed and shipped a company-wide interview-coordination automation spanning 180+ employees, 20+ hiring managers, and 5+ C-level executives across Germany, USA, and Tel-Aviv offices. 500+ interviews automated; scheduling friction reduced ~90%.

STACK

Languages & frameworks: Python, FastAPI, Flask, LangChain, LangGraph, SQLAlchemy (async), Alembic.

LLM & AI: Claude API, OpenAI, OpenRouter (multi-provider orchestration), Ollama (local), Cerebras, Deepgram, Minimax TTS; RAG with pgvector; multi-agent patterns; SSE streaming; YAML-driven prompt management.

Data & infra: PostgreSQL, pgvector, Docker, Railway, Gunicorn, Stripe, Google OAuth, Pandas, SQL.

PRIOR EXPERIENCE — TALENT ACQUISITION (DACH / EMEA)

Four years of full-cycle B2B SaaS / Fintech recruiting across engineering and GTM. Daily engineer- and stakeholder-facing work, requirements gathering, cross-functional coordination. These are foundational skills I now apply to shipping AI products.

- **Lucanet AG — Talent Partner (GTM, DACH/EMEA).** Oct 2023 – Feb 2026. 80%+ offer acceptance, ~120% of quarterly hiring targets; built Spain BDR hub and Italy GTM entry. Interviewed in EN, DE, ES, FR.
- **RapidAPI — Talent & Coordination Partner.** Jun 2022 – Aug 2023. Full-cycle tech recruiting (backend, frontend, DevOps, APIs) across three continents.
- **Volocopter — Talent Acquisition Specialist.** Jul 2021 – Jun 2022. Supported growth from ~200 to ~600 employees; aerospace, eVTOL, battery, software roles.
- **Microsoft & Honeywell — HR Generalist roles.** Mexico, 2016 – 2019. EMEA/LATAM/APAC HR operations.

EDUCATION & CERTIFICATIONS

- **Applied Generative AI Specialization — Purdue University.** Completed January 2026. LangChain, agents, Python automation, deployment.
- **MBA, International Business Administration — Karlsruhochschule, Karlsruhe.** 2019 – 2021.
- **BBA — Universidad Autónoma de San Luis Potosí, Mexico.** 2010 – 2016.
- **Hugging Face — AI Agents Fundamentals** (Feb 2025); Gradio Agents & MCP Hackathon (Jun 2025).

LANGUAGES

Spanish (native) · English (C2) · German (C1) · French (C1)